



UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas
in
Computer Architecture
(a.k.a. Machine Structures)

Memory (Mis)Management




cs61c.org



From ENIAC (1946) to EDSAC (1949)



- ENIAC: First Electronic General-Purpose Computer
- Needed 2-3 days to setup new program
- Programmed with patch cords and switches
 - At that time & before, "computer" mostly referred to people who did calculations
 - Mostly women! (See *Hidden Figures*, 2016)



- EDSAC: First General Stored-Program Computer
- Programs held as **numbers in memory**
 - Revolution! Program is also data!
- 35-bit binary **two's complement** words



05 Memory (Mis) Management (2)



What gets printed?

[Concept Check]

`sizeof()`: compile-time operator; gives size in bytes (of type or variable).

```
// for this exercise, assume
// shorts are 16b on a 64-bit architecture
void mystery(short arr[], int len) {
    printf("%d ", len);
    printf("%d\n", sizeof(arr));
}
```

```
int main() {
    short nums[] = {1, 2, 3, 99, 100};
    printf("%d ", sizeof(nums));
    mystery(nums, sizeof(nums)/sizeof(short));
    return 0;
}
```

- A. 10 5 10
- B. 10 5 8
- C. 80 5 80
- D. 80 5 40
- E. Other



05 Memory (Mis) Management (3)



What gets printed?

[Concept Check]

`sizeof()`: compile-time operator; gives size in bytes (of type or variable).

```
// for this exercise, assume
// shorts are 16b on a 64-bit architecture
void mystery(short arr[], int len) {
    printf("%d ", len);
    printf("%d\n", sizeof(arr));
}
```

```
int main() {
    short nums[] = {1, 2, 3, 99, 100};
    printf("%d ", sizeof(nums));
    mystery(nums, sizeof(nums)/sizeof(short));
    return 0;
}
```

Array has decayed
to a pointer

In array's declared scope,
total array size.

In array's declared scope,
elements in array.



05 Memory (Mis) Management (4)



Agenda

Memory Locations

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management




05 Memory (Mis) Management (5)



C Program Address Space

- A program's address space contains 4 regions:
 - **Stack**: local variables inside functions; grows downward
 - **Heap**: space requested via `malloc()`; resizes dynamically, grows upward
 - **Data (Static Data)**: variables declared outside main, does not grow or shrink
 - **Text (aka code)**: program executable loaded when program starts, does not change
- `0x00000000` chunk is unwriteable/unreadable so you that crash on NULL pointer access
- Programming in C requires knowing where objects are in memory, otherwise things don't work as expected.
 - By contrast, Java hides location of objects.



For now, OS somehow prevents accesses between stack and heap. (more later w/virtual memory)




05 Memory (Mis) Management (6)



Where are variables allocated?

- **Global**: If declared **outside** a function, allocated in **data (static)** storage.
- **Local**: If declared **inside** function, allocated on the **stack** and freed when function returns.
 - NB: `main()` is also a function.
- For both these memory types, the management is **automatic**.
 - You don't need to worry about deallocating when you are no longer using them.
 - But a variable **does not exist anymore** once a function ends!



05 Memory (Mis) Management (7)



Agenda

The Stack

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management




05 Memory (Mis) Management (8)





The Stack

- Every time a function is called, a new **stack frame** is allocated on the stack.
- A stack frame includes:
 - Return "instruction" address (who called me?)
 - Arguments*
 - Space for other local variables
- Stack frames contiguous blocks of memory; the **stack pointer** indicates the start of stack frames.
- When function ends, stack frame is tossed off the stack; frees memory for future stack frames.
- (more later when we cover details for a RISC-V processor architecture)



OS Memory (Mis) Management (9)

*see slide notes in pptx for technical elaboration




Garcia, Kuo



The Stack is Last In, First Out (LIFO)

```
int main ()
{ a(8); ... }
void a (int m)
{ b(1); }
void b (int n)
{ c(2); }
void c (int o)
{ d(3); }
void d (int p)
{ }
```



Stack grows down

The stack grows downward; a()'s local variables have lower byte addresses than main()'s, and so on.



OS Memory (Mis) Management (10)



Garcia, Kuo



Recall: Array Are Very Primitive [From last time]

- Array bounds are not checked during element access.
 - Consequence: We can accidentally access off the end of an array!
- An array is passed to a function as a pointer.
 - Consequence: The array size is lost! Be careful with `sizeof()`!
- Declared arrays are only allocated while the scope is valid.

```
char *foo() {
    char string[32]; ...;
    return string;
} is incorrect
```

Solution:
Dynamic memory allocation!
(more later now)



OS Memory (Mis) Management (11)



Garcia, Kuo




Passing Pointers into the Stack

It is fine to pass a pointer to stack space further down.

- I.e., pointers to addresses higher in the stack point to data in currently allocated stack frames.

```
void load_buf() { ... };
int main() {
    ... char buf[...];
    load_buf(buf, BUFLen);
    ... }
```



However, it is catastrophically bad to return a pointer to something in the stack!

- Memory will be overwritten when other functions are called!
- So your data would no longer exist...and writes can overwrite key pointers, causing crashes!

```
char *make_buf() {
    char buf[50];
    return buf;
}
void foo() { ... }
int main() {
    char *stackAddr = \
    make_buf();
    foo(stackAddr);
    ... }
```



OS Memory (Mis) Management (12)



Garcia, Kuo



Agenda

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management

The Heap



OS Memory (Mis) Management (13)



Garcia, Kuo



What is the Heap?

- The heap is dynamic memory – memory that can be allocated, resized, and freed during program runtime.
 - Useful for persistent memory across function calls.
 - But biggest source of pointer bugs, memory leaks, ...
 - Similar to Java **new** command allocates memory...but with key differences below.
- Huge** pool of mem (usually >> stack), but not allocated in contiguous order.
 - Back-to-back requests for heap memory *could* result in blocks very far apart
- In C, specify number of bytes of memory **explicitly** to allocate/deallocate item.
 - `malloc()`: Allocates raw, uninitialized memory from heap
 - `free()`: Frees memory on heap
 - `realloc()`: Resizes previously allocated heap blocks to new size
 - Unlike the stack, memory gets reused only when programmer **explicitly** cleans up



OS Memory (Mis) Management (14)



Garcia, Kuo



void *malloc(size_t n)

- Allocates a block of uninitialized memory:
 - `size_t n` is an unsigned integer type big enough to "count" memory bytes.
 - Returns `void *` pointer to block of memory on heap.
 - A return of NULL indicates no more memory (**always** check for it!!!)
- To allocate a struct:


```
typedef struct { ... } TreeNode;
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
```

Typecast casts return value from type (`void *`) to (`TreeNode *`) **sizeof(type)** gives size in bytes.
- To allocate an array of 20 ints:


```
int *ptr = (int *) malloc(20*sizeof(int));
if (ptr != NULL) { ... // always check NULL after
```

Many years ago ints used to be 16b. Now, 32b or 64b...

Assuming size of objects can lead to misleading, unportable code. Use `sizeof()`!



OS Memory (Mis) Management (15)



Garcia, Kuo



void free(void *ptr)

- Dynamically frees heap memory
 - `ptr` is a pointer containing an address originally returned by `malloc()`/`realloc()`.

```
int *ptr = (int *) malloc (sizeof(int)*20);
... free(ptr); // implicit typecast to (void *)
```

- When you free memory, be sure to pass the original address returned from `malloc()`. Otherwise, crash (or worse!)



OS Memory (Mis) Management (16)



Garcia, Kuo



void *realloc(void *ptr, size_t size)

- Resize a previously allocated block at ptr to a new size.
- Returns new address of the memory block.
 - In doing so, it may need to copy all data to a new location.
- realloc(NULL, size); // behaves like malloc
- realloc(ptr, 0); // behaves like free, deallocates heap block
- Remember: Always check for return NULL, which would mean you've run out of memory!

```
int *ip; ip = (int *) malloc(10*sizeof(int));
...
/* check for NULL */
ip = (int *) realloc(ip, 20*sizeof(int));
/* contents of first 10 elements retained */
...
/* check for NULL */
realloc(ip,0); /* equivalent to free(ip); */
```



OS Memory (Mis) Management (17)



Garcia, Kuo



Agenda

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management



OS Memory (Mis) Management (18)



Garcia, Kuo



Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```



OS Memory (Mis) Management (19)




Garcia, Kuo



Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



OS Memory (Mis) Management (20)



Garcia, Kuo




Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```



OS Memory (Mis) Management (21)




Garcia, Kuo



Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



OS Memory (Mis) Management (22)



Garcia, Kuo




Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```



OS Memory (Mis) Management (23)




Garcia, Kuo



Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



OS Memory (Mis) Management (24)




Garcia, Kuo



Linked List Example

```
#include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



Berkeley


Garcia, Kao



Linked List Example

```
#include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



Berkeley

Garcia, Kao



Agenda

When Memory Goes Bad

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management

Berkeley

OS Memory (Mis) Management (27)

Garcia, Kao



Working with Memory

Code, Static storage are easy:

- They never grow or shrink.


Stack space is also easy:

- Stack frames are created and destroyed in LIFO order.
- Just avoid “**dangling references**”: pointers to deallocated variables (e.g., from old stack frames).

⚠ Working with the heap is tricky:

- Memory can be allocated / deallocated at any time!
- “**Memory leak**”: If you forget to deallocate memory
- “**Use after free**”: If you use data after calling free
- “**Double free**”: If you call free 2x on same memory

Your program will eventually run out of memory



Possible crash or exploitable vulnerability

Berkeley

OS Memory (Mis) Management (28)

Garcia, Kao



Failure to free()

- The runtime does not check for the programmer’s failure to manage memory.
 - Memory is so performance-critical that there just isn’t time to do this.
 - Usual result: you corrupt the memory allocator’s internal structure, and you find out much later in a totally unrelated part of your code!
- **Memory leak**: Failure to free() allocated memory
 - ☹ Initial symptoms. Nothing...
 - Until you hit a critical point, memory leaks aren’t actually a problem
 - 📉 ...Later symptoms: performance drops off a cliff...
 - Memory hierarchy behavior tends to be great just up until it isn’t, then it hits several cliffs
 - 💀 ...and then your program is killed off!
 - Because the operating system (OS) says “no” when you ask for more memory

Berkeley

OS Memory (Mis) Management (29)

Garcia, Kao



Use after Free

- “**Dangling reference**”
 - When you keep using a pointer, even after it has been deallocated
- Reads after the free may be corrupted!
- If something else takes over that memory, your program will probably read the wrong information!
- Writes corrupt other data!
- Uh oh... Your program crashes later!

```
struct foo *f;
f = malloc(sizeof(struct foo));
...
free(f);
...
bar(f->a); // !!!
```

Berkeley

OS Memory (Mis) Management (30)

Garcia, Kao



Double-Free...

```
struct foo *f = (struct foo *) malloc(10*sizeof(struct foo));
...
free(f);
...
free(f); // !!!
```

- May cause either a use-after-free (because something else called malloc() and got that data) or corrupt heap data (because you are no longer freeing a pointer tracked by malloc)

Berkeley

OS Memory (Mis) Management (31)

Garcia, Kao



Forgetting realloc() Can Move Data

- “**Dangling reference**”
- Remember, when you realloc it can copy data to a different part of the heap.

```
int *nums;
nums = malloc(10*sizeof(int));
// forget to update nums
// on realloc call
realloc(nums, 20*sizeof(int));
// nums could now point
// to invalid memory,
// and we could have potentially
// lost a pointer to a new block
```

- Reads may be corrupted, and writes may corrupt other pieces of memory.

Berkeley

OS Memory (Mis) Management (32)

Garcia, Kao



Faulty Heap Management

[Concept Check]

How many memory management errors are in this code?

```

void free_mem_x() {
    int fnh[3];
    ...
    free(fnh);
}

void free_mem_y() {
    int *fum = malloc(4*sizeof(int));
    free(fum+1);
    ...
    free(fum);
    ...
    free(fum);
}

```

[for next time]

- A. 1
- B. 2
- C. 3
- D. 0
- E. Other



05 Memory (Mis) Management (33)



How many memory management errors are in this code?

1 0%

2 0%

3 0%

0 0%

Other 0%

Start the presentation to see live content. For screen share software, share the entire screen. Get help at policies.com/app



Faulty Heap Management

[Concept Check]

How many memory management errors are in this code?

```

void free_mem_x() {
    int fnh[3];
    ...
    free(fnh);
}

void free_mem_y() {
    int *fum = malloc(4*sizeof(int));
    free(fum+1);
    ...
    free(fum);
    ...
    free(fum);
}

```

free() on stack-allocated memory

free() on memory that isn't the pointer from malloc

Double free()

[for next time]

- A. 1
- B. 2
- C. 3
- D. 0
- E. Other



05 Memory (Mis) Management (35)



Valgrind to the rescue!!!

- Valgrind slows down your program by an order of magnitude, but...
 - It adds a tons of checks designed to catch most (but not all) memory errors
 - Memory leaks
 - Misuse of free
 - Writing over the end of arrays
- Tools like Valgrind are absolutely essential for debugging C code.

Check out Lab 02!



05 Memory (Mis) Management (36)



And in Conclusion...

- C has 3 pools of memory for variables:
 - Data: **global/static** variable storage, basically permanent
 - Stack: **local variable** storage, parameters, return address
 - Heap (**dynamic** storage): malloc() grabs space from here, free() returns it.
 - (4th memory pool: text, for the program executable itself)
- Heap data is biggest source of bugs in C code!



05 Memory (Mis) Management (37)



Agenda

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management

Implementing Memory Management

Material not tested. Recording:

<https://www.youtube.com/watch?v=5q5t5eWfn0Y>



05 Memory (Mis) Management (38)



Heap Management Requirements

- Want malloc() and free() to run quickly
- Want minimal memory overhead
- Want to avoid fragmentation*, when most of our free memory is in many small chunks
 - In this case, we might have many free bytes but not be able to satisfy a large request since the free bytes are not contiguous in memory.


* This is technically external fragmentation



05 Memory (Mis) Management (39)



Heap Management Example



05 Memory (Mis) Management (40)





K&R Malloc/Free Implementation

- From Section 8.7 of K&R
 - Code in the book uses some C language features we haven't discussed and is written in a very terse style; don't worry if you can't decipher the code
- Each block of memory is preceded by a header that has two fields:
 - size of the block, and
 - a pointer to the next block
- All free blocks are kept in a circular linked list.
- In an allocated block, the header's pointer field is unused.



K&R Malloc/Free Implementation

- `malloc()` searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- `free()` checks if the blocks adjacent to the freed block are also free.
 - If so, adjacent free blocks are merged (**coalesced**) into a single, larger free block.
 - Otherwise, freed block is just added to the free list.



Choosing a block in `malloc()`

- If there are multiple free blocks of memory that are big enough for some request, how do we choose which one to use?
 - **best-fit**: choose the smallest block that is big enough for the request.
 - **first-fit**: choose the first block we see that is big enough.
 - **next-fit**: like first-fit, but remember where we finished searching and resume searching from there.